


Broadcast Based Peer Review on Open Source Software Projects

Peter C. Rigby and Margaret-Anne Storey
University of Victoria
ICSE 2011



WTFs/minute

The ONLY valid measurement
of code quality: WTFs/minute



Why do peer review?

“The human eye has an almost infinite capacity for **not seeing what it does not want to see** [...] Programmers, if left to their own devices will ignore the most glaring errors in the output -- errors that **anyone else can see in an instant.**”

Weinberg

Background: OSS Peer Review

```
diff --git a/arch/mips/kernel/s
index 383aeb9..32a2561 100644
--- a/arch/mips/kernel/smp.c
+++ b/arch/mips/kernel/smp.c
@@ -193,6 +193,22 @@ void __dev
*/
```



[PATCH 1/1] Clocksource: Move the Hyper-V staging linux | x

- ★ **K. Y. Srinivasan** Move the Hyper-V clocksource driv
- ★ **Thomas Gleixner** Please do not use hard coded co
- ★ **john stultz** The mult/shift assignments can be drop
- ★ **Christoph Hellwig** > +#include <linux/version.h> Th
- ★ **KY Srinivasan** > -----Original Message----- > From: j
- ★ **KY Srinivasan** > devel@linuxdriverproject.org; john:
- ★ **KY Srinivasan** > devel@linuxdriverproject.org; tgix@
- ★ **Thomas Gleixner** to KY, gregkh, linux-l [show details](#)

What are the underlying **behaviours** and **mechanisms** that facilitate broadcast based review?

Method and Data

- Grounded Theory
 - Coded 460 review threads
 - Interviewed 9 core developers (top reviewers)

Style fix

```
+ flags |= UFS_ST_SUN;  
+ }  
+
```

Documentation/CodingStyle: if () {

reviewer comments
interleaved with code

Question
indicating
possible defect

```
+ switch (UFS_SB(sb)->s_flags & UFS_ST_MASK) {  
+ case UFS_ST_SUNOS:  
+     if (fs32_to_cpu(sb, usb3->fs_postblformat == UFS_42POSTBLFMT))
```

Really should be so?

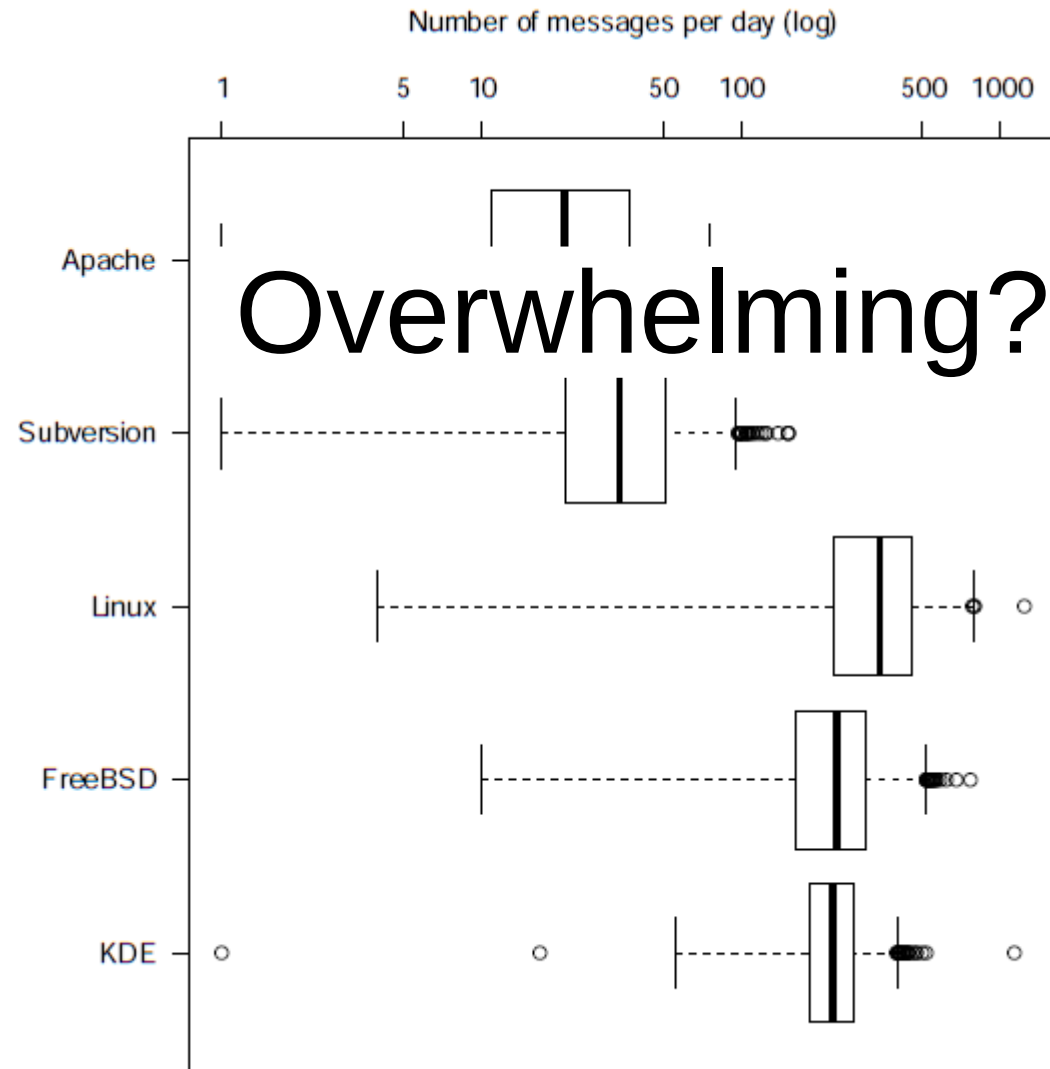

May be you mean:

fs32_to_cpu(sb, usb3->fs_postblformat) == UFS_42POSTBLFMT ?

We examine

- **Scalability** issues involved in broadcasting reviews in large projects,
- The techniques used to **find** patches for review,
 - Patch **structure** and norms,
- Review **outcomes**, stakeholders, interactions,
- The effect of too many **opinions** during a review, and
- The impact of **ignored** reviews

Mailing list traffic



A Day in Life for Greg KH

- **Total messages: 2067**
- 844 messages from mailing lists he skims
 - e.g., 595 lkml, 127 git
- 237 messages from 18 mailing lists that he read everything for
 - e.g., 51 meego-dev, 42 opensuse mailing lists
- 35 messages from kernel review requests
- 90 messages that are responses to reviews
- **2 personal emails**

Finding and Refinding



How do experts decide which contributions to review in a broadcast of general development discussion?



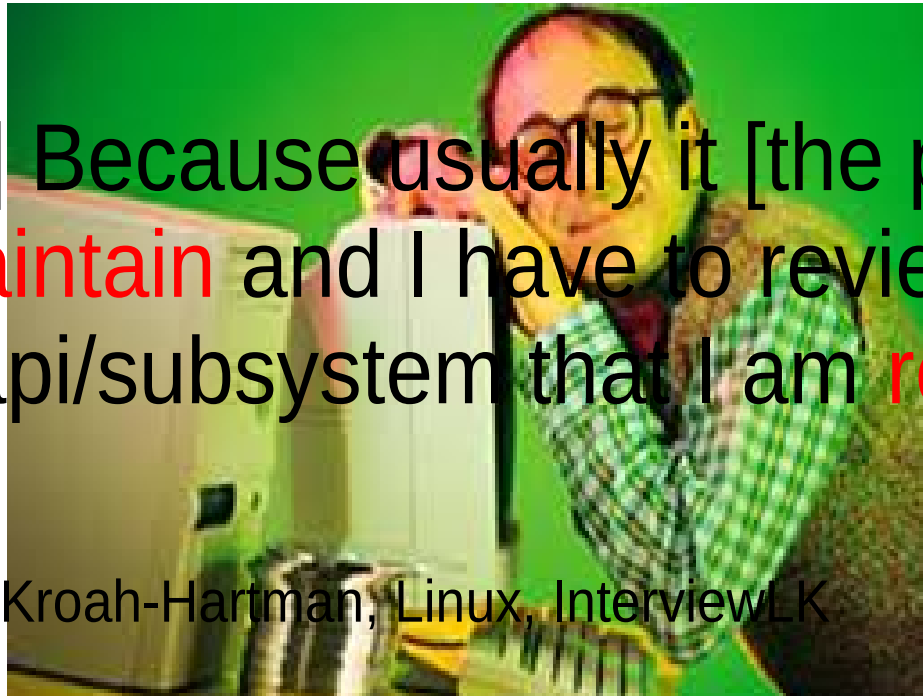
Reviewer Motivation

“Mostly [I review because of] **interest** or **experience** in the subject area [of the patch].”

Anonymous, FreeBSD, InterviewF2

“[I review] Because usually it [the patch] is for code I **maintain** and I have to review it. Or it uses an api/subsystem that I am **responsible** for.”

Kroah-Hartman, Linux, InterviewLK



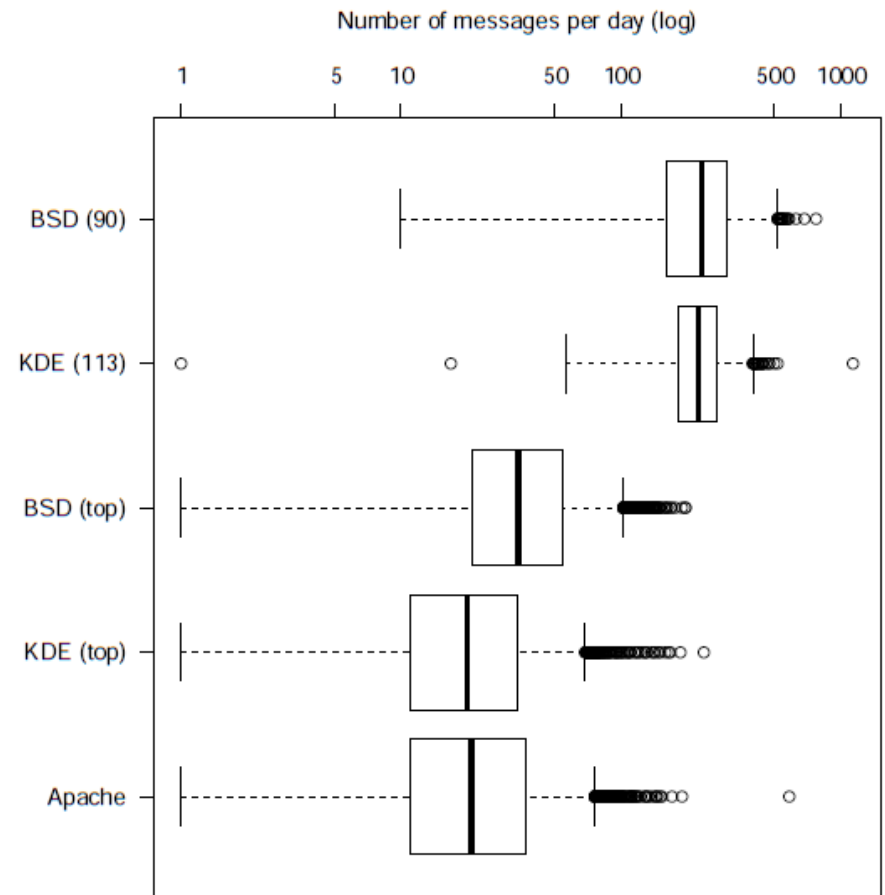
Scanning and Filtering



“Because of those filters, no, I do not feel overwhelmed [by the Linux Kernel Mailing List (LKML)]. On subsystem-specific mailing lists (like the linux-usb or linux-pci mailing lists), I read all emails.”

Multiple, Topic Specific Lists

- Multiple lists reduce traffic, but isolation?
- Few messages crosscut lists (5%)
- Many individuals post to multiple lists (25%)
- Subscribed to 5 and 30 lists
- Developers **Bridge** Lists



We examine

- **Scalability** issues involved in broadcasting reviews in large projects,
- The techniques used to find **patches** for review,
 - Patch **structure** and norms,
- Review **outcomes**, stakeholders, interactions,
- The effect of too many **opinions** during a review, and
- The impact of **ignored** reviews

Structure: Progressive detail

```
oom: oom_kill_process: fix the child_points logic
```

oom_kill_process() starts with victim_points == 0. This means that (most likely) any child has more points and can be killed erroneously.

Also, "children has a different mm" doesn't match the reality, we should check child->mm != t->mm. This check is not exactly correct if t->mm == NULL but this doesn't really matter, oom_kill_task() will kill them anyway.

Note: "Kill all processes sharing p->mm" in oom_kill_task() is wrong too.

Signed-off-by: Oleg Nesterov <oleg@redhat.com>

Signed-off-by: Linus Torvalds <torvalds@linux-foundation.org>

```
diff --git a/mm/oom_kill.c b/mm/oom_kill.c
index 7dcca55..b19c78e 100644
--- a/mm/oom_kill.c
+++ b/mm/oom_kill.c
```

“A good and descriptive **subject** line draws immediate attention.”

Structure: Progressive detail

```
oom: oom_kill_process: fix the child_points logic
```

```
oom_kill_process() starts with victim_points == 0. This means that  
(most likely) any child has more points and can be killed erroneously.
```

```
Also, "children has a different mm" doesn't match the reality, we should  
check child->mm != t->mm. This check is not exactly correct if t->mm ==  
NULL but this doesn't really matter, oom_kill_task() will kill them  
anyway.
```

```
Note: "Kill all processes sharing p->mm" in oom_kill_task() is wrong  
too.
```

```
Signed-off-by: Oleg Nesterov <oleg@redhat.com>
```

```
Signed-off-by: Linus Torvalds <torvalds@linux-foundation.org>
```

```
diff --git a/mm/oom_kill.c b/mm/oom_kill.c
```

```
index 7dcca55..b19c78e 100644
```

```
--- a/mm/oom_kill.c
```

```
+++ b/mm/oom_kill.c
```

Change log gives conceptual
understanding

Structure: Progressive detail

```
oom: oom_kill_process: fix the child_points logic
```

```
oom_kill_process() starts with victim_points == 0. This means that  
(most likely) any child has more points and can be killed erroneously.
```

```
Also, "children has a different mm" doesn't match the reality, we should  
check child->mm != t->mm. This check is not exactly correct if t->mm ==  
NULL but this doesn't really matter, oom_kill_task() will kill them  
anyway.
```

```
Note: "Kill all processes sharing p->mm" in oom_kill_task() is wrong  
too.
```

```
Signed-off-by: Oleg Nesterov <oleg@redhat.com>
```

```
Signed-off-by: Linus Torvalds <torvalds@linux-foundation.org>
```

```
diff --git a/mm/oom_kill.c b/mm/oom_kill.c  
index 7dcca55..b19c78e 100644  
--- a/mm/oom_kill.c  
+++ b/mm/oom_kill.c
```

The **diff** shows the file changes

Interleaved, Relevant History

Original diff

```
--- a/kernel/kexec.c
+++ b/kernel/kexec.c
@@ -1433,6 +1433,7 @@ module_init(crash_save_vmcoreinfo_init)
int kernel_kexec(void)
{
    int error = 0;
+   int locked;

    if (xchg(&kexec_lock, 1))
        return -EBUSY;
@@ -1498,7 +1499,8 @@ int kernel_kexec(void)
#endif

Unlock:
-   xchg(&kexec_lock, 0);
+   locked = xchg(&kexec_lock, 0);
+   BUG_ON(!locked);

    return error;
}
```

Snipped
content

```
On Wed, 13 Aug 2008, Huang Ying wrote:
>
> -   xchg(&kexec_lock, 0);
> +   locked = xchg(&kexec_lock, 0);
> +   BUG_ON(!locked);

Why do you want to do this at all?

And why do you implement your locks with xchg() in the
first place? That's total and utter crap.

Hint: we have _real_ locking primitives in the kernel.

Linus
```

Interleaved, Relevant History

Snipped
content

```
On Wed, 13 Aug 2008, Huang Ying wrote:
>
> -   xchg(&kexec_lock, 0);
> +   locked = xchg(&kexec_lock, 0);
> +   BUG_ON(!locked);
```

```
Why do you want to do this at all?
```

```
And why do you implement your locks with xchg() in the
first place? That's total and utter crap.
```

```
Hint: we have real locking primitives in the kernel.
```

Linus

Author
responds

```
Linus Torvalds <torvalds@linux-foundation.org> writes:
```

```
> On Wed, 13 Aug 2008, Huang Ying wrote:
```

```
>>
>> -   xchg(&kexec_lock, 0);
>> +   locked = xchg(&kexec_lock, 0);
>> +   BUG_ON(!locked);
```

```
> Why do you want to do this at all?
```

```
> And why do you implement your locks with xchg() in the first place? That's
> total and utter crap.
```

```
> Hint: we have _real_ locking primitives in the kernel.
```

```
This part certainly.
```

```
The way the code should work, and the way it has in the past is:
```

```
image = xchg(&kexec_image, NULL)
```

```
if (!image)
```

```
    return -EINVAL;
```

Refinding: Recipient Building

- Private email
 - Number of recipients doesn't change
- Mailing list
 - Number of recipients **increases** with the number of people who respond
 - Reviewers respond, refinding would be inefficient
 - Reply indicates interest, automatically added to **Cc list**, so
 - Message arrives in inbox as well as list

We examine

- **Scalability** issues involved in broadcasting reviews in large projects,
- The techniques used to **find** patches for review,
 - Patch **structure** and norms,
- Review **outcomes**, stakeholders, interactions,
- The effect of too many **opinions** during a review, and
- The impact of **ignored** reviews

Review Discussions

- Purely **technical** discussion
 - “Does the patch do what it claims to do and are there any flaws in the implementation?”, SW
- Scope, **politics**, and necessity of change arguments
 - Feature freeze
 - Misplaced or overly specialize fix
 - No significant improvement over existing code

Too many trivial opinions


- Parkinson's Law of triviality
- Lack of core group participation
- Exacerbated by large community?
- Impact of outsiders?



Outsider Involvement

	Apache	Subversion	FreeBSD	KDE
Percentage of reviews:				
Outsiders involved	30%	23%	13%	18%
Outsiders majority	5%	3%	2%	4%
“Outsiders’ influence” measure correlated with:				
Time for review	−.17	−.24	−.32	−.20
Total messages	−.53	−.59	−.60	−.38
Core-dev messages	−.76	−.77	−.77	−.63

Too Few Reviewers



“Lack of time can **postpone** the actual reviewing, but not really have an effect on whether to review or the **quality** of it.”

Faure, KDE, Interview, KF

We examined

- **Scalability** issues involved in broadcasting reviews in large projects,
- The techniques used to **find** patches for review,
 - Patch **structure** and norms,
- Review **outcomes**, stakeholders, interactions,
- The effect of too many **opinions** during a review, and
- The impact of **ignored** reviews

Takeaways

- Invested co-developers review because of **interest and obligation**
 - Simple tools, techniques, and social norms to find reviews
- **Postpone** review rather than rush
 - Ignore incompetent, biased authors
- Asynchronous reviews facilitate **group discussion** and roles that accord with varied skills
- **Politicized** pointless discussion happen infrequently
- Scaling through **multiple lists** and explicit requests

Threats to Credibility

- **Fit**: send results to interviewees
 - 5 interviewees responded, all positive
- **Triangulation**: multiple methods and data
 - Archival and interview data
 - Grounded theory and quantification of parameters
- **Generalizability**
 - Successful, mature OSS projects
 - 4 of 5 infrastructure
- Publicly available data

Theory: Parameters of Review

1. Early, frequent reviews
2. of small, independent, complete contributions
3. that, despite being asynchronously broadcast to a large group of stakeholders, are reviewed by a small group of self-selected experts,
4. resulting in an efficient and effective peer review technique.